Project Green — Business Plan

Full document assembled from sections 1–6. Individual section files (01–06) remain unchanged in this folder.

Part 1 — Problem, Opportunity & Why Now

1. The Problem: A Structural Crisis, Not a Cyclical One

The US accounting industry is not in a temporary downturn. It is in a permanent workforce and operational transformation. Firms cannot hire their way out of it.

The Talent Crisis

- **Retirement wave:** Nearly three-quarters of CPAs are at or approaching retirement age. The pipeline is not recovering: CPA exam candidates have declined sharply, and the 150-hour education requirement discourages new entrants.
- **Knowledge drain:** When senior staff leave, decades of institutional knowledge—how the firm handles edge cases, client history, unwritten procedures—leave with them. Junior staff onboarding takes months because so much is undocumented.
- **Partner bottleneck:** Partners spend a large share of their day answering questions their team should be able to answer: policy lookups, procedural guidance, "how did we handle this last time?" The most expensive resource in the firm is consumed by repetitive, low-leverage work.

The Operational Impact

- **Inconsistent answers** erode client trust and create liability when different staff give different guidance.
- **Slow responsiveness:** Clients ask "Why did my profit drop?" or "What's my AR balance?" and wait days for an answer that could be instant if the firm's data were queryable in one place.
- **No single view:** Firms run on QuickBooks, payroll systems, tax software, document stores, and spreadsheets. No product today connects all of them into a unified intelligence layer. Questions that span systems go unanswered or require manual assembly.

Why Current Substitutes Fail

Substitute	Why it fails
Ask a senior	Bottleneck, interruption, inconsistent answers; doesn't scale.
Search QBO / SharePoint	Slow, limited to one system, often fruitless.
Manual dashboards	Stale by the time they're built; don't answer ad hoc questions.
Paste data into ChatGPT	No live connection to firm data; confidentiality risk; not firm-specific.
Hire more people	The talent does not exist in the volume needed.

The industry has the budget—$145B+ in the US—but not the supply of human expertise to execute. AI is not optional; it is the only scalable response to the workforce gap.

---

2. The Opportunity: An Underserved Segment and an Unfilled Gap

Market Size and Structure

- **80,000+ accounting firms** in the US; the vast majority are small to mid-sized (roughly 78,000 firms in the 2–50 person range).
- **AI accounting automation** is a multi-billion-dollar market growing at roughly 40% CAGR; adoption in accounting jumped from single digits to over 40% in a single year. Firms are actively looking for solutions.
- **Serviceable segment:** 10–50 person firms doing tax and bookkeeping, on cloud GL (e.g. QuickBooks Online), in reachable metros, with managing partners who feel capacity and knowledge pain. This is a multi-billion-dollar serviceable market, with early-year targets in the tens of firms and $100K+ ACV potential per firm.

The Gap No One Fills

- **Multi-system intelligence:** No product connects a firm's GL, payroll, tax software, and documents into one queryable layer. Point solutions automate one workflow (AP, close, bookkeeping); ERPs stay inside one ecosystem. When a partner asks "How is Client X performing across all our services?", no tool can answer today.
- **Institutional knowledge for small firms:** When a senior leaves a 15-person firm, their knowledge walks out. Enterprise knowledge tools are too heavy; generic AI doesn't know the firm's procedures or data. No productized solution captures and serves that knowledge for firms this size.
- **The 78,000 small firms:** The Big Four and large vendors build for the top of the market. Startups serve small firms with narrow point solutions. The segment that needs firm-wide intelligence—same core needs as larger firms, but smaller budget and no internal IT—is underserved.

The opportunity is to be the **operating intelligence layer** for these firms: connect what they already use, ingest their data and documents, and deliver a firm-specific AI that answers questions accurately, with citations, without replacing their existing systems.

---

3. Why Now: Five Converging Forces

The problem and the opportunity are not new in the abstract. What is new is that **now** the technology, market, and structure of the industry align to make a solution viable and urgent.

1. Agentic AI Is Production-Ready

Multi-step, autonomous AI agents that plan, decide, act, and verify across workflows are no longer research projects. 2025–2026 is when this capability became viable for production. The leap from "chatbot" to "agent that uses your data and tools" is what makes a firm-wide assistant possible. The technology is ready.

2. Private Equity Is Reshaping the Industry

Private equity is rolling up accounting firms at scale. A significant share of the largest US firms are or will be PE-backed. These owners need operational clarity, standardized processes, margin optimization, and exit-ready metrics. When rates normalize and exits ramp up, firms that can demonstrate operational intelligence will command premium valuations. The demand for a "process twin" and clear operational data is accelerating.

3. The CPA Retirement Wave Is Biting Now

The retirement wave is no longer a future scenario; it is current. Firms cannot replace departing expertise through hiring alone. The urgency to automate, preserve knowledge, and reduce dependence on a few key people is immediate. Solutions that capture institutional knowledge and scale answers are directly aligned with this timing.

4. A Fragmented Market of 80,000+ Firms

Unlike enterprise software with a handful of dominant buyers, accounting is a large, fragmented market with tens of thousands of firms that share the same core needs: better use of data, faster answers, preserved knowledge. A

productized solution that works for one 15-person firm can scale across many without customizing for each. The structure of the market supports a platform play.

5. Big Four Invest in AI but Not in This Segment

The Big Four have committed billions to AI and are building digital twins and proprietary tools—for their own practices and for the largest clients. They are not productizing for the 78,000 smaller firms. That creates awareness and demand ("we need something like that") without supply. The gap between what the Big Four advertise and what small and mid-sized firms can access is the opening.

---

4. The Timing Window

The window is open but finite.

- **2025–2026:** Adoption inflection. Firms are evaluating AI; early movers can establish data moats and switching costs.
- **2027–2028:** Market consolidation. Incumbents (e.g. Intuit, Microsoft) will ship more capable AI; winners will be those already in production with paying customers.
- **2029+:** Mature market. Late entrants will face entrenched players with deep integration and accumulated data.

The imperative: **a working product and paying customers by the end of 2026, and scaling into 2027.** Project Green's first alpha customer (February 2026) validates that the problem is real and that a firm will commit to using the product; the next step is delivery, then conversion and replication.

Summary

- **Problem:** Structural CPA workforce and knowledge crisis; partners bottlenecked; no single intelligence layer across firm systems; current substitutes don't scale.
- **Opportunity:** 80,000+ firms, a multi-billion-dollar serviceable market, and a clear gap: firm-wide, multi-system intelligence and institutional knowledge capture for small and mid-sized firms.
- **Why now:** Agentic AI is viable; PE roll-ups and exit dynamics create demand for operational intelligence; retirement-driven urgency is current; market fragmentation supports productized scale; Big Four create demand they don't serve for this segment.
- **Timing:** Act in 2025–2026 to capture the adoption window; delay cedes the window to incumbents and faster-moving competitors.

Part 2 — Solution & Product

1. The Solution in Brief

Project Green builds an **AI-powered operating intelligence layer** for professional service firms. We do not replace the firm's existing systems—QuickBooks, payroll, tax software, document stores—we connect to them, aggregate the data, and deliver a **firm-specific AI assistant** that answers questions in plain English with accurate, cited answers drawn from the firm's own data and documents.

The product is an **intelligent overlay**: it makes the accountant's entire firm queryable, reduces partner interruption, captures institutional knowledge, and—over time—enables client-facing explanation and process-level visibility. We are not an ERP, not a general-purpose chatbot, and not a compliance tool. We are the layer that sits on top of what firms already use.

---

2. Two-Layer Model: Platform + Services

Layer 1: Productized Core (The Platform)

The platform we build, own, and deliver:

Component	Role
Multi-agent orchestration	Routes questions to the right specialist; coordinates tool calls; formats answers for the user. Built on LangGraph with an orchestrator and domain specialists (firm-level, client-level, document).
Data lakehouse	Single data layer that aggregates all firm systems. Medallion architecture: raw ingestion (bronze) → normalized schema (silver) → business-ready views (gold). Same agent logic works whether the source is QuickBooks, Xero, or payroll—data is normalized at the silver layer.
RAG knowledge base	Vector store for firm SOPs, policies, memos, and past work. Enables "ask the firm's knowledge" without manual documentation. Improves over time as more documents and Q&A are ingested.
Security and access control	Multi-tenant isolation (every query scoped by firm); conversation and role-aware access; audit trails for production.
Deployment adapters	Thin presentation layer: standalone web UI today; Microsoft Teams and Google Workspace later. Core logic is platform-agnostic via the Model Context Protocol (MCP).

Layer 2: Professional Services (How We Deploy)

Firms do not self-serve this product. We deploy it with them:

- **Assessment** — Review the firm's systems, data quality, and readiness (lead-gen or flat fee).
- **Twin setup** — Connect data sources, configure agents, load firm knowledge (one-time setup fee).
- **Ongoing tuning** — Refine behavior, add use cases, update the knowledge base (monthly retainer).

The pitch: **"We build your firm's AI twin."** This is a services-plus-software model (similar in spirit to Palantir or Salesforce): the services revenue supports the build; the platform and accumulated data create the long-term moat.

---

3. Three Product Archetypes (Sequenced)

We deliver value in three stages. Each archetype builds on the previous one's data and infrastructure.

Archetype 1: Internal Staff Copilot (MVP — 2026)

What it is: A private, firm-specific AI used only by firm staff. Trained on the firm's financial data (via the lakehouse) and documents (SOPs, policies, memos via RAG). Staff ask in plain English; the AI answers with citations. "What's our policy on capitalizing expenses over $500?" "What was Client X's AR balance last month?" "How do we handle month-end close for clients on cash basis?"

Why first: Internal-only, so no client liability. Highest immediate value—reduces partner interruption (the most expensive resource), speeds onboarding, and ensures consistent answers. Proves the platform before we add client-facing or process-modeling complexity.

Success looks like: The firm uses it daily without being prompted; partners report measurable time back; new hires get answers from the AI instead of interrupting seniors.

Archetype 2: Client-Facing Explanation Engine (V2 — 2027)

What it is: A client-safe AI that **explains** financial data—not advice, only explanation. Example: "Your operating expenses increased 12% this month, driven mainly by contractor payments on Jan 15 and Jan 22. [Source: GL entries #1842, #1856]." Guardrails keep the AI from crossing into advisory or recommendations.

Why second: Requires live financial data and firm-approved language; liability and tone need to be designed and tested. Builds on the same lakehouse and agent stack; adds guardrails, citation standards, and client-facing UX.

Archetype 3: Process Twin (Long-Term — 2028+)

What it is: A digital model of how work flows through the firm. Questions like: "Where do we lose time in month-end close?" "What happens if we add two associates?" "Which clients are most profitable per hour?" Simulation and optimization for capacity, staffing, and client mix.

Why third: Highest value but highest complexity—process data, simulation, and change-management. Natural buyers are PE-backed firms that need operational clarity and exit readiness. We get there after Archetypes 1 and 2 are proven and the data foundation is in place.

---

4. How It Works: Architecture Summary

- **User** asks a question (today: web chat; later: Teams or other surfaces).
- **Orchestrator** interprets the question and routes to the right specialist: firm-level financials, per-client financials, or document/knowledge.
- **Specialist agents** call tools (via MCP) against the lakehouse or the RAG store. They do not call vendor APIs directly; they query the normalized lakehouse so one integration (e.g. QuickBooks) can be swapped or extended without rewriting agent logic.
- **Answers** are grounded in retrieved data and documents, with citations. A verifier step checks answers against source data where appropriate. We do not generate financial numbers without a source.
- **Data flow:** Vendor systems (QuickBooks Online first; then Xero, Gusto, etc.) feed the lakehouse. Firm documents feed the RAG pipeline. Both are isolated per firm (multi-tenant).

What we are NOT building: We are not an ERP (we don't replace QuickBooks or any system of record). We are not a general-purpose AI (we are firm-specific). We are not a compliance or audit tool (operational questions and intelligence, not regulatory filings). We are not a generic data warehouse product—the lakehouse is infrastructure that powers the AI, not a standalone offering.

---

5. Current State and Roadmap to Production

Current State (as of 2026)

The platform is built and measurable:

Dimension	Status
Multi-agent system	Orchestrator + firm, client, and document specialists; 8 tool modules (ledger, analytics, staff, time/billing, scheduling, tax, payroll, documents).
Correctness	120+ automated tests with structured evaluation framework; iterating toward production accuracy.

Data layer	DuckDB/DuckLake medallion architecture in place; QuickBooks Online integration in progress.
RAG	ChromaDB-based pipeline in progress for document Q&A.
Interface	Streamlit chat UI with message history and tool-call visibility.
Quality assurance	120+ automated tests; evaluation harness for regression.

Development uses a local LLM (e.g. Qwen3:32B via Ollama) for zero marginal inference cost; production can use the same or a cloud model. The architecture is model-agnostic.

Phased Roadmap (Summary)

Phase	Focus	Outcome
Phase 0 (now – Q2 2026)	Foundation	QBO + RAG live; system answers 10 representative questions correctly with real data and documents.
Phase 1 (Q3–Q4 2026)	Staff Copilot MVP	One beta firm (alpha customer) onboarded; daily use; feedback loop; path to paid conversion.
Phase 2 (Q1–Q2 2027)	Multi-vendor, multi-surface	Xero and Gusto integrations; Teams surface; second firm; onboarding repeatable.
Phase 3 (Q3 2027+)	Explanation Engine and scale	Archetype 2 live with guardrails; production infra (e.g. cloud lakehouse, vector DB); 5–10 firms.
Phase 4 (2028+)	Process Twin	Process modeling and simulation; PE-focused packaging; cross-firm benchmarking; optional vertical expansion (law, healthcare).

The first alpha customer (February 2026) is the forcing function for Phase 1: we are committed to delivering a production-ready Staff Copilot they can rely on, then converting to paid and replicating.

---

6. Strategic Choices That Shape the Product

- **Platform-agnostic core via MCP:** Integrations and presentation (Teams, web, etc.) are adapters. We avoid locking into one vendor before we need to; Microsoft is the likely first surface for the target segment.
- **DuckLake for MVP:** Full lakehouse behavior (ACID, time travel, schema evolution) without cloud dependency. We move to cloud when multi-tenant scale and compliance require it.
- **QuickBooks first:** Most common GL among target firms; mature SDK and sandbox. Fastest path to real accounting data through the system.
- **Narrow ICP for launch:** 10–50 person firms, tax + bookkeeping, QuickBooks Online, reachable metro, managing partner who feels capacity and knowledge pain. One segment, one use case, prove it before widening.

Summary

- **Solution:** An AI operating intelligence layer that connects to a firm's existing systems and delivers a firm-specific assistant with cited, accurate answers.

- **Model:** Productized platform (orchestration, lakehouse, RAG, security, adapters) plus professional services (assessment, setup, tuning). "We build your firm's AI twin."
- **Archetypes:** (1) Internal Staff Copilot (MVP, 2026)—internal only, highest immediate value; (2) Client-Facing Explanation Engine (2027)—explanation, not advice; (3) Process Twin (2028+)—operational model and simulation, PE-oriented.
- **Architecture:** Orchestrator + specialists, MCP-based tools, medallion lakehouse, RAG. We do not replace ERPs; we overlay them.
- **Status:** Multi-agent system and evaluation in place; QBO and RAG in progress; alpha customer secured; roadmap is Phase 0 → 1 (beta with one firm) → 2 (multi-vendor, second firm) → 3 (Explanation Engine, scale) → 4 (Process Twin).

Part 3 — Market Size & Competition

1. Market Size

US Accounting Industry Structure

The US accounting industry comprises **80,000+ firms**. The vast majority are small to mid-sized:

Segment	Firms	Employees per Firm	Annual Revenue per Firm	Tech Spend (est.)
Large (Top 100)	~100	500–50,000+	$100M–$20B+	$5M–$500M+
Mid-market	~2,000	50–500	$5M–$100M	$250K–$5M
Small (target)	**~78,000**	**2–50**	**$200K–$5M**	**$5K–$100K**

Total industry revenue is on the order of **$145B+** (US). The AI accounting automation segment is multi-billion today and projected to grow at roughly **40% CAGR** over the next decade. Adoption in accounting jumped from single digits to over 40% in a single year; 77% of firms plan to increase AI investment. The market is in early-stage transformation with real budget and intent to spend.

TAM / SAM / SOM

Two complementary views:

Spend potential (all firms that could use AI operational intelligence):

Metric	Definition	Value
TAM	All US accounting firms × reasonable annual AI/operational-intelligence spend	~$8B/year (80,000 firms × ~$100K potential spend)
SAM	10–50 person firms on cloud GL (QBO/Xero), open to AI, reachable	~~$2B/year (~~20,000 firms × ~$100K)
SOM (Year 1–2)	Target metros, warm-networkable, active capacity/knowledge pain	~$5M/year (e.g. 50 firms × ~$100K)

Our revenue opportunity at target ACV (~$36K–$48K):

Metric	Firms	Blended ACV	Annual Value

TAM	~15,000 (mid-tier, our fit)	~$36K	~$540M/year
SAM	~6,000 (cloud, AI-open, reachable)	~$38K	~$228M/year
SOM (Year 3)	50–100	~$40K	$2M–$4M/year
SOM (Year 5)	200–500	~$44K	$8.8M–$22M/year

The takeaway: even 50 firms at target ACV is a meaningful early-stage business; the market is deep enough to support **$100M+ ARR** at scale. We do not need to win the whole market to build a large, defensible company.

Target Segment (ICP)

We focus first on **10–50 person accounting firms** that:

- Do both **tax and bookkeeping** (year-round workflow, not tax-only).
- Use **QuickBooks Online** (or equivalent cloud GL) so we can integrate without replacing systems.
- Are in **reachable metros** for founder-led sales and support.
- Have a **managing partner who feels** capacity constraints, knowledge loss, or partner bottleneck.

This segment is large enough to feel the pain of the structural crisis (Section 1), has budget for a solution (IT and managed services spend in the $5K–$100K range), and is underserved by both enterprise vendors (who focus up-market) and point-solution startups (who automate one workflow, not the whole firm). Sub-segments such as **outsourced/virtual firms** (cloud-native, tech-forward, margin pressure) may prove especially receptive and are a natural extension of the same ICP.

---

2. Market Dynamics (Forces That Shape Demand)

Five forces reinforce demand for what we build:

1. **PE consolidation** — PE is rolling up accounting firms; one-third of the top 30 US firms may be PE-backed. Portfolio firms need operational clarity, standardization, and margin data. When rates normalize and exits ramp, firms with operational intelligence command premium multiples. The Process Twin (Archetype 3) is the PE product; PE operating partners are a future channel.
2. **CPA workforce crisis** — ~75% of CPAs at retirement age; firms cannot hire their way out. The Staff Copilot (Archetype 1) directly addresses knowledge capture and partner interruption.
3. **AI adoption inflection** — Adoption jumped 4.5× in one year; 82% of early adopters report positive ROI in year one. The window is open; firms are actively evaluating AI.
4. **Big Four investment** — $9.5B committed to AI; they build for their own practices and large clients, not for the 78,000 smaller firms. They create awareness and demand without serving this segment.
5. **Rising client expectations** — Clients expect faster, clearer answers ("Why did my profit drop?"). Advisory is growing faster than compliance. The Explanation Engine (Archetype 2) and a more responsive firm (enabled by Archetype 1) meet this demand.

---

3. Competitive Landscape

Competitive Matrix

Category	Examples	Approach	Threat to Project Green
Enterprise ERP	Oracle Fusion, SAP Joule, Dynamics	AI on top of large ERP	**Low** — wrong segment, wrong price point

	365		
Mid-market ERP	Acumatica, NetSuite	AI features in cloud ERP	**Low–Medium** — closer to our segment but ERP-centric, not overlay
SMB ERP	Business Central, Odoo	Copilot-style features for small business	**Medium** — right segment, but generic rather than firm-specific
AI accounting startups	Truewind, Vic.ai, Numeric, Gravis AI	Point solutions or vertical AI	**Medium–High** — closest competitors; narrow scope or different deployment model
Big Four	Deloitte, PwC, KPMG, EY	Proprietary tools, consulting	**Low** — different segment; they validate the market and don't productize for SMB
Horizontal AI	ChatGPT, Claude, Gemini	General-purpose assistants	**Medium** — firms may try these first; we win when the question requires firm-specific data (e.g. "What was Client X's AR balance?")

Startup Competitors (Closest Comparables)

Company	Focus	Our differentiation
Truewind	AI bookkeeping automation	Point solution for bookkeeping; we span firm-level financials, client-level, documents, and (long-term) process.
Vic.ai	Accounts payable automation	AP-specific; we model the entire firm and answer cross-system questions.
Numeric	Month-end close acceleration	Close-specific; we address the full operational lifecycle and firm knowledge.
Gravis AI	AI ERP for service businesses	Closest to our vision; they position as ERP replacement. We are overlay —connect to existing systems (QBO, Xero, etc.) and do not require migration.

Where Project Green Wins

Dimension	Typical competitors	Project Green
Data scope	Single system or one workflow	All firm systems aggregated in a lakehouse; one query can span GL, payroll, documents.
Agent model	Single copilot or task-specific bot	Multi-agent orchestration with domain specialists (firm, client, document) and routing.
Deployment	Replace your software or stay inside their app	Overlay on whatever the firm already uses; no rip-and-replace.
User model	One interface for everyone	Architecture supports role-aware routing and responses (partner vs. associate vs. client).
Long-term value	Task automation ("do this faster")	Operational intelligence ("understand your entire firm"); data accumulates and switching cost grows.

What we are not: We are not an ERP. We are not a general-purpose AI. We win when the question requires **live, firm-specific data** or **firm-specific knowledge**—the moment a user asks "What was account 1000's balance last month?" or "What's our policy on X?" using their data, generic AI and single-workflow tools fail. That gap is the core of our positioning.

---

4. Competitive Moat

What makes Project Green defensible over time:

1. **Firm-specific data accumulation** — Every day the product runs, it gets better at that firm's terminology, questions, and processes. Migrating to a competitor means losing that accumulated intelligence. Vertical SaaS switching costs for mid-market customers are often in the hundreds of thousands of dollars; we build the same dynamic.
2. **Cross-firm network effects** — With anonymization, we can offer benchmarks: "Firms like yours typically close in X days." Each new firm improves the value for every existing firm. The first 10–20 firms don't get much here; at 50+ firms, benchmarking becomes a real moat.
3. **Integration depth** — Deep integrations to QuickBooks, Xero, Gusto, ADP, ProConnect, and document stores are expensive to build and maintain. New entrants must replicate that connectivity.
4. **Services relationship** — The "we build your firm's AI twin" wrapper creates trust and a human relationship. Relationship-based switching cost compounds with data-based switching cost.
5. **PE channel** — If PE operating partners recommend us to portfolio firms, we get a distribution channel that scales with PE activity: each new acquisition is a potential deployment.

---

5. Expansion Markets (Later)

The same architecture—lakehouse, multi-agent orchestration, RAG, overlay deployment—applies to other knowledge-intensive professional service firms. Not part of the initial plan but relevant to long-term TAM:

Market	Size / growth	Fit
Law firms (small)	$4B AI legal tech, ~31% CAGR	High — same knowledge-capture and multi-system pattern
Healthcare practices	$36B AI medical billing by 2034	Medium–high — alpha customer can validate
Fractional CFO / FAO	$55B market, ~8% CAGR	Very high — identical "firm data + documents" architecture
Financial advisory (RIAs)	~15,000 firms	Medium — similar client-serving and data pattern

We do not pursue these before proving the model in accounting. Vertical expansion is a Phase 4+ option once the Process Twin and PE channel are in play.

---

6. Timing Window (Recap)

The market window is open but finite (see also Part 1):

- **2025–2026:** Adoption inflection. Firms are evaluating AI; early movers can establish data moats and switching costs.

- **2027–2028:** Consolidation. Incumbents (e.g. Intuit, Microsoft) will ship more capable AI; winners will be those already in production with paying customers.
- **2029+:** Mature market. Late entrants will face entrenched competitors with deep integration and accumulated data.

We must have a **working product and paying customers by end of 2026** and be **scaling by 2027** to capture the window. Our first alpha customer (February 2026) is the first step; the next is delivery, conversion, and replication.

Summary

- **Market size:** 80,000+ US accounting firms; ~$8B TAM (spend potential), ~$540M TAM at our ACV; SAM in the hundreds of millions; SOM Year 1–2 in the single-digit millions (tens of firms). Market is large enough for $100M+ ARR at scale.
- **Target segment:** 10–50 person firms, tax + bookkeeping, cloud GL (QBO), reachable, managing partner who feels capacity/knowledge pain. Underserved by both enterprise and point-solution startups.
- **Competition:** Enterprise and mid-market ERP (low threat); AI startups (medium–high, narrow scope or ERP-replacement); Big Four (low, validate demand); horizontal AI (medium—we win on firm-specific data and knowledge).
- **Moat:** Firm-specific data accumulation, cross-firm benchmarking, integration depth, services relationship, PE channel potential.
- **Timing:** Capture the adoption window with product and paying customers by end of 2026; scale in 2027 or cede advantage to incumbents and faster-moving competitors.

Part 4 — Business Model, GTM & Sales

1. Business Model Overview

Project Green operates a **services-plus-software** model: a productized platform (the AI twin) wrapped in professional services (assessment, setup, tuning). We do not sell pure SaaS; firms buy an outcome—"we build your firm's AI twin"—delivered through a combination of software and human delivery. Revenue comes from one-time setup, recurring platform access, and optional ongoing services. This structure matches how firms already buy managed IT, fractional CFO, and practice management: they expect a relationship and implementation support, not self-serve tooling, when the product touches their financial data and operations.

---

2. Revenue Streams and Pricing

Revenue Streams

Stream	Type	Timing	Description
Twin Setup	One-time	At onboarding	Connect data sources (QBO, etc.), configure agents, load firm knowledge (SOPs, policies).
Platform Fee	Recurring monthly	Ongoing	Core access to the AI twin; includes a defined number of users.
Tuning Retainer	Recurring monthly	Optional, ongoing	Monthly check-ins, new use cases, knowledge base updates, optimization.

Expansion Projects	One-time or recurring	As needed	Add Archetype 2 (Explanation Engine), new integrations (Xero, Gusto), Process Twin modules; project- or module-priced.

Pricing Structure

Setup + platform + optional retainer:

Component	Range	Notes
Firm Assessment	Free	Lead generation; demonstrates value and builds trust.
Twin Setup	$5,000–$15,000 one-time	Covers integration, configuration, knowledge loading.
Monthly Platform	$2,000–$5,000/month	Varies by firm size and tier; includes up to N users.
Ongoing Tuning Retainer	$1,000–$3,000/month (optional)	Monthly check-ins, new use cases, knowledge updates.

Tiered ACV by firm size (representative):

Tier	Firm Revenue	ACV (approx.)	Monthly Platform
Starter	$2M–$10M	$18K–$30K	$1,500–$2,500/mo
Growth	$10M–$25M	$36K–$48K	$3,000–$4,000/mo
Scale	$25M–$50M	$48K–$72K	$4,000–$6,000/mo

Target ACV for first 5 firms: $30,000–$70,000 per firm (blended). Full-year target ACV at scale: ~$40K (validated against vertical SaaS medians, accounting firm IT budgets, managed services, practice management software, and ROI from recovered partner time).

Pricing Validation

Pricing is anchored to existing budget categories (managed IT, fractional CFO, practice management) and to value (e.g. recovery of 20% of one FTE's time). Validation plan: test in 10+ firm conversations; use first customer(s) as beta/discounted engagements in exchange for case studies and feedback; refine based on willingness-to-pay and perceived value, not cost.

---

3. Unit Economics (Summary)

Cost to serve one firm:

- **LLM inference:** ~$5–$50/month depending on model mix and query volume (local or low-cost cloud models).
- **Infrastructure (MVP):** ~$0–$70/month (local or single VPS). At production scale: ~$150–$670/month per firm (hosting, lakehouse, vector DB, monitoring).
- **Human time:** 20–40 hours one-time setup; 6–12 hours/month ongoing (tuning, support). This is the main capacity constraint for a small team.

Gross margin: Software-like. Phase 1: ~97–98% (revenue $2K–$3.5K/month per firm, cost ~$50). Phase 2–3: ~86–94% (revenue $3.5K–$8K/month, cost ~$300–$500). Local-first and efficient model routing keep inference cost low versus cloud-heavy competitors (55–70% gross margin), enabling either competitive pricing or reinvestment in the services layer.

Revenue expansion: Net Revenue Retention above 120% is expected from: (1) Archetype 1 → Archetype 2 expansion, (2) additional integrations and users within the firm, (3) Process Twin and benchmarking as paid upgrades. Vertical SaaS benchmarks support this range. Detailed projections are in Part 5 (Financials).

---

4. Ideal Customer Profile and Disqualifiers

Primary ICP (Phase 1–2):

Attribute	Criteria	Rationale
Firm size	10–50 employees	Large enough to feel pain, small enough to decide quickly.
Services	Tax + bookkeeping / business management	Same data, maximum value from one integration.
General ledger	QuickBooks Online	First integration; most common in segment.
Geography	Within reach of founder (driving distance or strong relationship)	Founder-led sales and support for first 3–5 firms.
Pain	Capacity constraints; too much work, not enough people; knowledge loss	CPA shortage and partner bottleneck as entry wedge.
Decision maker	Managing partner or ops partner who feels the pain	Single decision maker who can say yes.
Technology stance	Open to AI but not yet committed to a solution	Easier to win in a new category than to displace an incumbent.

Secondary ICP (Phase 3+): 50–200 person firms; PE-backed (operating partner as champion); full-service (tax, audit, advisory); need for operational standardization across acquired practices.

Disqualifiers (for now): Solo practitioners (too small); desktop-only software (no cloud API); no documented processes (nothing for RAG); firms in active M&A transition (too distracted).

---

5. Positioning and Messaging

Core message: "*Ask this before you ask a human.*"

Outcome-based, not technology-based. It speaks to the pain (interruptions, capacity, slow answers) without requiring the buyer to understand AI, agents, or lakehouses.

By persona:

Persona	Pain	Our message

Managing Partner	"I spend half my day on questions my team should know."	"Your AI twin knows what you know. Your team asks it first."
Ops Partner	"Every new hire takes months to get up to speed."	"New hires get answers on day one. Your SOPs, your policies, your way."
Senior Accountant	"I get interrupted constantly with the same questions."	"The AI handles repeatable questions. You handle the ones that matter."
Junior Associate	"I don't know who to ask or where to find the answer."	"Ask the AI. It knows your firm's procedures and points you to the right person when it doesn't."

What we do not say: We do not say "AI will replace your staff." We do not lead with technical jargon (LangGraph, RAG, etc.). We do not position as an ERP. We do say "your firm gets smarter every day" (aspiration, growth).

Dual mode for lead gen vs. product: For lead generation (cold/warm outreach, content, conferences), we lead with "AI digital transformation" and a consultative, problem-first frame. For demos and pilot proposals, we switch to product positioning: "Ask this before you ask a human" and show the product on representative data.

---

6. Sales Motion

Phase 1: Founder-Led Sales (First 5 Firms)

Mike (and Owen where relevant) personally sell, onboard, and support. No dedicated sales team or paid funnel. Reliance on personal relationships and warm introductions.

Process:

1. **Identify** — Map networks for accounting firm connections (personal contacts, LinkedIn, local CPA society, ProAdvisor networks).
2. **Warm introduction** — Get introduced to managing partner or ops partner. No cold outreach for the first 5.
3. **Free assessment** — "Let us look at your firm's systems and processes. We'll show you where AI can help." Lead-gen offer; builds trust.
4. **Demo** — Show the product on representative accounting data (e.g. QBO sandbox). "This is what your team's experience would look like."
5. **Pilot proposal** — "Connect your QuickBooks, load your SOPs, run a 30-day pilot with 3–5 staff. Here's what we'll measure."
6. **Pilot** — 30-day trial with active support, weekly check-ins, rapid iteration.
7. **Convert** — Convert to paid annual contract when value is proven.

Alpha customer (secured Feb 2026): First "customer" is an alpha partner at 99% discount in exchange for alpha partnership. This validates that the "discount for alpha" tactic works, that founder-led sales works, and that the pain is real. Pricing validation is deferred until conversion to paid or engagement of the second customer. The flywheel starts with this engagement: case study, reference, feedback, and warm intros to peers.

Phase 2: Referral-Driven (Firms 6–20)

- Every satisfied firm generates referrals to 2–3 peers.
- Case studies fuel conference talks, webinars, and thought leadership.
- QBO ProAdvisor and accounting tech consultant networks as referral channels.
- Content and conferences (e.g. state CPA societies, Scaling New Heights, Intuit Connect) drive inbound interest.

Phase 3: Channel Sales (20+ Firms)

- **PE operating partners** — One PE relationship can mean 5–20 portfolio accounting practices. PE needs operational intelligence, standardization, and margin data; we become the recommended platform.
- **Accounting technology consultants** — Referral partners who include us in tech stack recommendations.
- **Industry associations** — Partnerships and visibility through AICPA, state societies, and accounting technology communities.

---

7. The Flywheel

Each new firm increases the value of the product and the efficiency of sales:

- **More data** → better benchmarks and patterns (anonymized) for all firms.
- **More questions and usage** → better agent behavior and training signal.
- **More use cases** → broader product coverage and expansion revenue.
- **More case studies** → stronger proof and easier conversion of the next firm.

One firm delivers value → case study and thought leadership → conference talks and referrals → next firm → expanded usage and cross-firm data → product more valuable for every firm. The flywheel is the same whether we stay bootstrap or raise capital; the only variable is how fast we spin it.

---

8. Partnerships and Content

Partnerships (build in Phase 1–2): QBO ProAdvisors (access to accounting firm network), local CPA societies (speaking, credibility), accounting technology consultants (referrals). **Phase 2–3:** PE operating partners, Intuit/QuickBooks (co-marketing, API), Microsoft (M365 Agents SDK, distribution).

Content and thought leadership: LinkedIn 2–3×/week (trends, operations, build updates); blog 2×/month (deep dives on AI and accounting); conference talks quarterly; podcast appearances monthly; case study per customer. Conferences: AICPA ENGAGE, Intuit Connect, Xerocon, Scaling New Heights, AccountingWEB Live, state CPA society meetings. See also lead-gen target lists for organizations and directories.

---

9. Customer #2 and Beyond

With the alpha engagement in place, the next customer should test a different variable: closer to standard pricing (e.g. 50% discount rather than 99%), ideally a referral from the alpha customer, same ICP (10–50 person, QBO, capacity-constrained), and willing to provide a case study. Each subsequent customer de-risks repeatability of onboarding and pricing.

Summary

- **Business model:** Services-plus-software. Revenue from setup (one-time), platform fee (recurring), optional tuning retainer, and expansion projects. Target ACV $30K–$70K for first 5 firms; ~$40K at scale.
- **Unit economics:** High gross margin (97–98% Phase 1, 86–94% Phase 2–3); main constraint is human time (6–12 hrs/month per firm). NRR >120% from archetype expansion, integrations, and modules.
- **ICP:** 10–50 person firms, tax + bookkeeping, QBO, reachable, managing partner who feels capacity/knowledge pain. Disqualify solo, desktop-only, no processes, active M&A.
- **Positioning:** "Ask this before you ask a human." Persona-specific messages; no fear-based or technical positioning.
- **Sales:** Phase 1 founder-led (warm intro → assessment → demo → pilot → convert); Phase 2 referral and content; Phase 3 PE and consultant channels.

- **Alpha customer:** Validates founder-led motion and "discount for alpha"; flywheel and pricing validation start with conversion and customer #2.

Part 5 — Financial Projections & Unit Economics

1. Unit Economics

Cost to Serve One Firm

LLM inference (monthly): ~$5–$50 per firm depending on model mix (e.g. GPT-4o-mini or Gemini Flash for most queries; premium models only when needed) and query volume. Assumptions: on the order of 50 queries/day per firm, ~500 input + ~200 output tokens per query (multi-turn), 22 working days/month. Strategy: model routing, caching of common answers, and RAG precision keep cost low. Local models in development keep marginal cost near zero.

Infrastructure (monthly):

Component	Phase 0–1 (MVP)	Phase 2–3 (Production)
Compute / hosting	$0 (local) or ~$20 (VPS)	$50–$200 (cloud VM/containers)
Data lakehouse	$0 (DuckLake local)	$50–$300 (e.g. Databricks, S3+compute)
Vector database	$0 (ChromaDB local)	$25–$70 (Qdrant, Pinecone)
LLM inference	$0 or $5–$50	$5–$50
Monitoring	$0	$20–$50
Total per firm	**~$0–$70/month**	**~$150–$670/month**

Human time (capacity constraint): 20–40 hours one-time setup per firm; 6–12 hours/month ongoing (tuning, support). At 6–12 hours/month per firm and ~~50% of a two-person team allocated to customer work, the team can support **~~10–15 firms** before a dedicated customer-success or delivery hire is required. This is the binding constraint in Phase 1–2, not infrastructure cost.

Gross Margin by Phase

Phase	Monthly Cost per Firm	Monthly Revenue per Firm	Gross Margin
Phase 1 (MVP)	~$50	$2,000–$3,500	**97–98%**
Phase 2–3 (Production)	~$300–$500	$3,500–$8,000	**86–94%**

Margins are software-like. The main cost at scale is human delivery (setup and tuning), which is captured in the pricing (setup fee and optional retainer). Local-first and efficient model routing keep LLM and infra cost low versus cloud-heavy competitors (often 55–70% gross margin), enabling competitive pricing or reinvestment in the services layer.

Revenue Expansion and NRR

Net Revenue Retention above **120%** is expected from: (1) Archetype 1 → Archetype 2 expansion, (2) additional integrations (Xero, Gusto) and users within the firm, (3) Process Twin and benchmarking as paid upgrades. Vertical

SaaS leaders average 120%+ NRR; our expansion levers align with that pattern.

LTV / CAC (Directional)

Metric	Value	Basis
Average customer lifetime	5+ years	Vertical SaaS retention; switching costs (data moat, integration depth)
Target annual churn	< 10%	Vertical SaaS leaders often < 8%
Blended annual gross profit per firm	~$40K	Weighted by tier mix (Starter/Growth/Scale)
Customer LTV	$200K+	5 years × ~$40K gross profit
CAC (founder-led)	$5K–$10K	Time, travel, demo costs
CAC (scaled)	$15K–$25K	Sales, marketing, events
LTV/CAC	8–20×	Above the 3× minimum threshold

---

2. Revenue Projections

Scenario Conventions

- **Conservative:** Slower customer adds; lower ACV; minimal new hires.
- **Moderate (plan):** Alpha converts to paid; 2–5 new firms per quarter; ACV ramps with tier mix.
- **Aggressive:** Faster adds and higher ACV; used for upside cases only.

All figures are directional. Alpha customer is currently at 99% discount; first "revenue" in projections assumes conversion to paid or a second paying customer.

Moderate Scenario (Plan)

Year	Firms	Avg ACV	ARR	Cumulative Revenue
2026 (H2)	3	$30K	$90K	~$60K
2027	15	$38K	$570K	~$460K
2028	40	$42K	$1.68M	~$1.59M
2029	80	$44K	$3.52M	~$4.20M
2030	140	$46K	$6.44M	~$9.18M
2031	220	$48K	$10.56M	~$18M

Quarterly Build (Moderate) — First 18 Months

Quarter	Firms	Setup + Recurring	Quarterly Revenue	Cumulative
Q3 2026	1	1 beta/alpha convert	~$11K–$15K	~$11K–$15K

Q4 2026	2–3	New setups + MRR	~$17K–$42K	~$28K–$57K
Q1 2027	4–6	New setups + MRR	~$34K–$95K	~$62K–$152K
Q2 2027	6–10	New setups + MRR	~$46K–$155K	~$108K–$307K
Q3 2027	10–15	New setups + MRR	~$80K–$265K	~$188K–$572K
Q4 2027	15–20	New setups + MRR	~$115K–$355K	~$303K–$927K

Year 1 (mid-2026 to mid-2027) revenue: ~$100K–$150K (conservative to moderate).
Year 2 run rate (end 2027): ~$360K (conservative) to ~$1.4M (moderate).

Conservative Scenario (Slower Ramp)

Quarter	Firms	Monthly Recurring	Quarterly Revenue	Cumulative
Q3 2026	1 (beta)	$2,000	~$11,000	~$11,000
Q4 2026	2	$4,000	~$17,000	~$28,000
Q1 2027	4	$8,000	~$34,000	~$62,000
Q2 2027	6	$12,000	~$46,000	~$108,000
Q3 2027	10	$20,000	~$80,000	~$188,000
Q4 2027	15	$30,000	~$115,000	~$303,000

Year 1: ~$108K. **Year 2 run rate:** ~$360K ARR.

Key Revenue Milestones (Moderate Plan)

Milestone	Target Timing
First dollar of revenue	Q3 2026 (alpha convert or second customer)
First $100K ARR	Q2 2027
$1M ARR	Q1 2028
$5M ARR	Q2 2030
$10M ARR	Q4 2031

---

3. Cost Structure

Fixed Costs (Monthly)

Category	Phase 0–1	Phase 2 (1–20 firms)	Phase 3 (20+ firms)
Tools (Cursor, collaboration, domain, etc.)	~$300–$400	~$300–$500	~$500–$1,000
Cloud / infrastructure (shared)	$0–$50	$200–$500	$500–$2,000

LLM API (shared)	$0–$50	$50–$200	$200–$1,000
Total fixed	**~$300–$400**	**~$600–$1,200**	**~$1,200–$4,000**

Fixed costs scale with firm count and product maturity (e.g. production lakehouse, monitoring). Until multi-tenant production, a large share of cost is variable per firm.

Variable Costs (Per Firm)

See "Cost to Serve" above: **~$50/month (Phase 1)** to **~$300–$500/month (Phase 2–3)**. Variable cost is dominated by infrastructure and LLM at scale; human time is the capacity constraint, not the marginal cost of one more firm once the team is in place.

People Costs (When They Enter)

Role	When	Approx. Annual Cost
Mike (no salary from company)	Stage 1 (friends round)	$0 (personal expenses covered by W-2 + practice income)
Owen (part-time retention)	Stage 1 (friends round)	$12K ($1,000/month)
Mike (salary begins)	Stage 2 (pre-seed)	$72K ($6,000/month)
Owen (expanded to 20+ hrs)	Stage 2 (pre-seed)	$48K ($4,000/month)
First hire (engineer or customer success)	When firm count exceeds ~10–15 or seed funding	$60K–$100K

The founder takes no salary during Stage 1. Personal expenses (~~$4K/month) are covered by the W-2 salary and ML Financial Group income ($30K/year). The founder already self-funds the entire business at ~$750-$1,200/month — Owen at $115/week (authorized to double to 10 hrs/week at $230/week), AI tools (~~$200/month), subscriptions (~$75/month). Owen's line item in the use of funds formalizes his existing compensation under the entity. The founder has a ~$60K personal safety net ($40K 401K + $20K investments) separate from business capital. Stage 1 raise ($50K) is pure growth capital — legal foundation, formalized co-founder compensation, conferences, and deployment hosting. See Part 6 for the staged "ask" and use of funds.

---

4. Key Financial Milestones and Decision Gates

Milestone	Target Date	Significance
First dollar of revenue	Q3 2026	Alpha convert to paid or second paying customer; validates willingness to pay.
$10K MRR	Q4 2026	3–5 firms at $2K–$3.5K/month; recurring base established.
Infrastructure self-funding	Q1 2027	Revenue covers cloud, API, and tooling costs.
$50K MRR	Q3 2027	~10–15 firms; capacity limit of two-person team; trigger to consider first hire or capital.

$100K MRR	Q4 2027	Sustainable base; decide bootstrap vs. raise for acceleration.

---

5. Funding Context (Summary)

Capital efficiency: The founder is self-funding the entire business today at ~~$750-$1,200/month — including co-founder compensation (Owen at $115/week, authorized to double to 10 hrs/week), AI tools (~~$200/month), and subscriptions (~~$75/month). Personal expenses (~~$4K/month) are covered by the W-2 salary and ML Financial Group income ($30K/year). A personal safety net of ~$60K ($40K 401K + $20K investments) exists separately from business capital. **No external capital is required for the business to continue operating.** The raise is pure growth capital — every dollar goes to legal foundation, growth activities, and acceleration.

Staged external capital path: The fundraise is structured in three stages, each gated by milestones:

1. **Friends Round ($50K, SAFE at $3M–$3.5M cap):** Entity formation, legal, first customer deployment, path to first revenue. Led by a strategic accounting firm investor (prospective first customer). Closeable in weeks with one check. Founder takes no salary; entire amount is growth capital.
2. **Pre-Seed ($300K–$500K, SAFE at $6M–$8M cap):** Team expansion, engineering, go-to-market execution. Raised after 1–3 paying customers prove the model. Milestone: 5–10 paying firms, $200K+ ARR.
3. **Seed ($1M–$2M, priced round at $10M–$20M):** Full team, PE channel, Archetype 2 production. Raised after significant ARR and repeatable sales motion.

Bootstrap path remains viable: Development funded by sweat equity and personal savings; first customers fund infra and tools. The staged approach is not survival-driven—it is an acceleration strategy that preserves equity by raising at progressively higher valuations. Any specific "ask" (amounts, instruments, use of funds) is in Part 6.

---

6. Valuation Implications (Reference)

At **10× revenue** (median for vertical SaaS): $100K ARR ≈ $1M valuation; $500K ≈ $5M; $1M ≈ $10M; $5M ≈ $50M; $10M ≈ $100M. Early-stage (seed) multiples often 15–25× ARR; later stages trend toward 8–12×. Used for context in investor discussions and optionality (bootstrap vs. raise); not a commitment to a specific valuation.

---

7. Sensitivity and Risk (Financial)

- **Slower acquisition:** If customers close in 9 months instead of 3, revenue and cash flow shift right; runway and milestone dates extend. Bootstrap remains viable with a longer path to $100K ARR.
- **Lower ACV:** If blended ACV is ~$24K instead of ~$40K (e.g. 40% lower), ARR and valuation drop proportionally; unit economics stay profitable (e.g. 82%+ gross margin on Starter tier). Business remains viable; growth and valuation are slower.
- **Higher burn:** Adding a full-time hire before revenue or funding shortens runway; recommended only if customer acquisition is ahead of plan or capital is secured.

Summary

- **Unit economics:** ~$50/month cost per firm (Phase 1), ~$300–$500 (Phase 2–3); gross margin 97–98% / 86–94%. Human time (6–12 hrs/month per firm) caps a two-person team at ~10–15 firms. NRR >120% from expansion.
- **Revenue (moderate plan):** 2026 H2 ~$90K ARR (3 firms); 2027 ~$570K (15 firms); 2028 ~$1.68M (40 firms); 2030 ~$6.44M (140 firms); 2031 ~$10.56M (220 firms). Key milestones: first revenue Q3 2026, $100K ARR Q2 2027, $1M ARR Q1 2028.

- **Cost structure:** Fixed ~$300–$400/month (Phase 0–1) to ~$1,200–$4,000 (Phase 3); variable ~$50–$500 per firm per month. First hire when firm count or funding justifies it.
- **Funding:** Bootstrap viable to ~10+ firms; raise when opportunity exceeds capacity or competition demands speed. Valuation context: 10× revenue median; 15–25× at seed.
- **Sensitivity:** Slower acquisition or lower ACV extends timeline or reduces valuation; unit economics remain profitable. Avoid adding fixed cost (e.g. full-time hire) before revenue or funding unless acquisition is ahead of plan.

Part 6 — Traction, Team, Ask & Risks

1. Traction

Milestones Achieved

Date	Milestone	Significance
2026-02-22	**First alpha test customer secured**	Agreement signed: 99% discount in exchange for alpha partnership. Project Green has a committed customer to deliver to. Validates market interest and shifts the project from prototyping to production delivery.

The alpha engagement confirms that (1) the "discount for alpha" tactic works, (2) founder-led sales works for this segment, and (3) the pain is real—a firm was willing to invest time in an early product.

Product and Platform Status (as of 2026)

Dimension	Status
Multi-agent system	Orchestrator + firm, client, and document specialists; 8 tool modules (ledger, analytics, staff, time/billing, scheduling, tax, payroll, documents).
Correctness	120+ automated tests with structured evaluation framework; iterating toward production accuracy.
Evaluation and quality	120+ automated tests; benchmarking framework for regression.
Data layer	DuckDB/DuckLake medallion architecture; QuickBooks Online integration in progress.
RAG	ChromaDB-based pipeline in progress for document Q&A.
Interface	Streamlit chat UI with message history and tool-call visibility.

The product is demonstrable today. Remaining Phase 0 work (QBO live, RAG production-ready) is required for the alpha customer to use it in daily operations, not for first conversations.

2026 Objectives and Status

Objective	Success criteria	Status
Ship a beta product	Working Staff Copilot answering real questions for one firm; firm uses it daily without prompting	**Committed** — alpha customer secured 2/22; production delivery underway

Develop go-to-market plan	Documented ICP, messaging, pricing, sales motion; messaging tested with 10+ conversations	In progress — GTM docs complete; alpha validates "99% discount for alpha"
Land Owen in full-time employment	Skills development and capital generation	Unchanged (external to Project Green)
Land first paying customer	Revenue from one firm using the product; customer retains after 90 days	**On track** — alpha is the path to conversion; paid engagement follows successful alpha
Build the growth flywheel	Case study from first customer; 3+ warm leads; pipeline of 5+ qualified firms	Dependent on alpha success

Next Steps (Post–Alpha)

1. **Deliver to alpha** — Harden security, error handling, and deployment; connect customer's QBO and documents; meet success criteria (operational system, trustworthy answers, 3–5 staff using it, path to paid).
2. **Convert alpha to paid** — Transition from 99% discount to paid arrangement when value is proven.
3. **Customer #2** — Ideally a referral from alpha; test pricing closer to standard (e.g. 50% discount); same ICP; case study and reference.
4. **Pipeline** — Use alpha success for warm introductions to 2–3 additional firms; begin testing messaging with 10+ firm conversations.

---

2. Team

Founders

Mike Leavenworth — CEO & Founder (75% equity)
18+ years in accounting and finance. Founder, ML Financial Group (accounting practice since 2017; active client of 9 years). Reviewed financial reporting across 60+ accountants for high-net-worth clients. Career path: internal audit at Fortune 500 → category management → founded own practice → building AI for the industry he knows. Stanford Lean Startup; Institute of Internal Auditors Board Member. Domain expertise that cannot be hired—the founder is the customer.

Owen Peterson — Co-Founder & CTO (25% equity, 4-year vesting)
AI/ML engineering and research. 10+ year collaborative relationship with the founding team through the open-source community. Transitions to expanded role with funding; developing into CTO position. As we scale, we plan to bring on an experienced CTO through equity dilution shared across founders. This hire would work collaboratively with Owen and the CEO in a knowledge-sharing arrangement designed to accelerate technical leadership across the team.

Advisory Board (6 formal advisors, 0.25% equity each)

Domain	Qualification
Cybersecurity & Compliance	Sr Cybersecurity Analyst, BMO. US Navy veteran. 20+ years. NIST 800-53, PCI-DSS, vendor risk management.
AI & Engineering	Deep AI and coding expertise. Entrepreneurial builder with hands-on depth in modern AI tooling.
Marketing & GTM	Marketing strategy and go-to-market. Critical for scaling.

Operations & Problem-Solving	Versatile, committed, action-oriented.
Business Judgment	Grounded, practical reality-check perspective.
Legal & Dealmaking	Active search; finalizing. Seeking counsel with startup, corporate, and transaction expertise to advise on entity formation, fundraising, and strategic partnerships.

Informal Network

A senior software architect (30+ years, principal-level at a major technology company) serves as an informal mentor and provides technical guidance through a shared community. Not a formal advisor due to conflict-of-interest constraints.

---

3. The Ask (Fundraising Path)

If Project Green raises external capital, the fundraise is structured in stages to minimize dilution and maximize proof points before scaling. The business plan is also executable on a bootstrap path (see Part 5); the ask is for the fundraising scenario.

Stage 1: Friends Round

Term	Detail
Instrument	SAFE (Simple Agreement for Future Equity)
Amount	$50,000–$75,000
Valuation cap	$3,000,000–$3,500,000 (post-money)
Discount	None (standard post-money SAFE)
Lead investor	Strategic accounting firm (prospective first deployment customer)

Use of funds:

Category	Amount	Purpose
Legal & entity formation	$15K	Attorney review + opinion, C-Corp, IP assignment, E&O insurance
Owen compensation (12 mo)	$12K	$1K/month co-founder retention
Alpha deployment & hosting	$5K	Production hosting for first customer(s)
Conferences (2 events)	$6K	Travel + registration for customer acquisition
Sales materials	$3K	One-pager, demo environment, basic website
Reserve	$9K	Buffer for unexpected costs

Note: the founder takes no salary from the company during Stage 1. Personal expenses are covered by the W-2 and practice income. The founder already self-funds the entire business (including Owen's compensation) at ~$750-$1,200/month. Owen's line item formalizes existing compensation under the entity. This raise is pure growth capital.

Milestone: Entity formed, legally clear, alpha deployed, 1–3 paying customers, first revenue.

Stage 2: Pre-Seed

Term	Detail
Instrument	SAFE
Amount	$300,000–$500,000
Valuation cap	$6,000,000–$8,000,000 (post-money, justified by revenue)
Prerequisite	1–3 paying customers, recurring revenue, deployed product

Use of funds (12–18 month runway):

Category	Allocation	Purpose
Team expansion	40%	Owen to 20+ hrs/week, founder salary increase, first contractor hires
Engineering	25%	Multi-vendor integrations (Xero, Gusto), security hardening, production infrastructure
Go-to-market	20%	Sales execution, conference circuit, content engine, customer acquisition
Infrastructure & working capital	15%	Cloud hosting at scale, compliance prep, contingency

Milestone: 5–10 paying firms, $200K+ ARR, positioning for seed round at $10M–$20M valuation.

Why Staged

The founder already self-funds the entire business (including co-founder compensation) at ~$1,200/month and takes no salary. The friends round ($50K at $3.5M cap = ~1.4% dilution) is pure growth capital, closeable in weeks with one check. Stage 2 happens with revenue proof at a higher valuation ($6M–$8M cap), preserving more equity. Total pre-seed dilution across both stages is ~7.1% for $450K raised — compared to ~8.9% for $400K in a single raise at $4.5M cap.

---

4. Risk Factors and Mitigations

We disclose material risks and how we address them. Probability and impact are directional.

Competitive Risks

Risk	Severity	Mitigation

Big Four build SMB product	Low prob., high impact	Move fast; build firm-specific data moats and integration depth. Big Four move slowly; target 50+ firms before any downmarket move.
ERP incumbents (e.g. Intuit) add agent features	High prob., medium impact	Position as overlay connecting all systems; Intuit cannot be the multi-system layer (e.g. won't integrate Xero).
Well-funded startup wins segment	Medium	Speed and focus; small team ships fast. Most startups are point solutions; we are multi-domain.
Horizontal AI (ChatGPT, Claude) "good enough"	Medium	Generic AI cannot access firm data or enforce access control. Demo the gap: "What was account 1000's balance?" requires firm data.

Technical Risks

Risk	Severity	Mitigation
LLM hallucination on financial data	High	Ground every answer in retrieved data; mandatory citations; verifier node; regression testing with 120+ automated tests.
Data quality from vendor APIs	Medium	Silver-layer data quality checks; flag issues to firm as a service; start with firms with reasonably clean books.
Integration complexity	Medium	One integration first (QBO); MCP isolates integrations; budget 2× time per integration.

Regulatory and Liability Risks

Risk	Severity	Mitigation
AI liability for accounting advice	High impact	Archetype 1 internal-only; Archetype 2 explains, never advises; guardrails and terms; citations for verification.
Data privacy regulations	Medium	Data in firm-controlled environments where possible; encryption; multi-tenant isolation; clear data agreements.

Market and Execution Risks

Risk	Severity	Mitigation
Firms adopt AI slower than expected	Medium	Free assessment and 30-day pilot reduce commitment; internal-only Archetype 1 is lowest-friction deployment.
Two-person team capacity	High	Ruthless prioritization; one firm, one use case at MVP; automate (CI/CD, testing); first hire at 10–15 firms.
Owen's split focus	Medium	Planned; modular architecture allows parallel work; Owen's employment can fund infra.
Long sales cycles	Medium	Founder-led sales and free assessment; 30-day pilot reduces perceived risk.

Price sensitivity at target ACV	Medium	Tiered pricing; $40K below vertical SaaS median; ROI anchored to recovered FTE time.

Alpha-Specific Risks

Risk	Severity	Mitigation
Product readiness gap	High	Security hardening and error handling first; production readiness checklist; deploy with Docker.
Alpha expectations exceed delivery	Medium	Clear scope; transparency on what the system can and cannot do; structured feedback and triage.
Reputation risk from poor alpha	Medium	Treat alpha as production-grade delivery; 24-hour response on issues; avoid scope creep.

Other

Risk	Severity	Mitigation
Data security concerns	High	Advisor with NIST/PCI-DSS expertise; security-first architecture; SOC 2 preparation at seed.
Employment / IP constraints	Medium	Governed by California law with strong employee protections (Labor Code § 2870, B&P Code § 16600); IP developed independently on own time/equipment; attorney review in progress; if current employer invests (Stage 1 strategy), resolves by mutual agreement.

Summary

- **Traction:** First alpha customer secured Feb 2026 (99% discount for alpha partnership). Product: 8 tool modules, 120+ automated tests, evaluation framework, chat UI; QBO and RAG in progress. 2026 objectives: ship beta to alpha, first paying customer via conversion, build flywheel. Next: deliver to alpha, convert to paid, land customer #2.
- **Team:** Mike (CEO, 75%) — 18+ years accounting/finance, domain founder. Owen Peterson (CTO, 25%, 4-year vesting) — AI/ML engineering, expanded role with funding; experienced CTO hire planned as a collaborative addition. Six formal advisors (0.25% each): security, AI/engineering, marketing/GTM, operations, business judgment, legal/dealmaking (finalizing). Informal technical mentor.
- **Ask (fundraising path):** Staged raise — Stage 1 friends round $50K–$75K SAFE at $3M–$3.5M cap (entity, legal, first deployment, first revenue); Stage 2 pre-seed $300K–$500K SAFE at $6M–$8M cap after revenue (team, engineering, GTM). Total pre-seed dilution ~7.9% for $475K. Seed at $10M–$20M after significant ARR. Bootstrap path remains viable (Part 5).
- **Risks:** Competitive (Big Four, ERP, startups, horizontal AI); technical (LLM reliability, data quality, integration); regulatory (AI liability, privacy); market (adoption speed); execution (capacity, Owen split focus, sales cycles, price sensitivity); alpha (readiness, expectations, reputation); security and employment. Mitigations in place for each; early warning indicators and prioritization (e.g. security and alpha delivery first) reduce downside.